MICROSOFT HEALTH AI · MARCH 2026

One bet. *Trust* is the next frontier in health AI.

This project turns *trust* in consumer health AI from a vague UX concern into a measurable model behavior by evaluating how systems respond to real patient questions.

I validated this two ways: benchmarked four models on trust and scraped 836 patient reviews across consumer health apps. Here is what I found and what I would build at Microsoft.

8%	836	30 ADVERSARIAL PROMPTS ACROSS 4 MODELS, 5	0

DROP IN PUBLIC TRUST IN HEALTH AI, 2025 TO 2026

APP STORE REVIEWS SCRAPED ACROSS 4 PLATFORMS

DIMENSIONS

EXISTING BENCHMARKS FOR CONSUMER HEALTH AI TRUST

THE ARGUMENT

Benchmarks like MedQA measure accuracy. Trust behaviors are less systematically evaluated.

Clinical accuracy tells you if the model is right. In consumer AI, that is not enough. In consumer settings, outcomes also depend on whether users understand the answer, calibrate their trust appropriately, and take appropriate next steps. A model can be correct and still fail through overconfidence, unclear escalation, or poor emotional alignment.

The behaviors that precondition trust are measurable: Calibrated uncertainty, appropriate escalation, transparent limitation, and emotional attunement. Starke et al. 2025 identifies these across health AI contexts. Teams have yet to operationalize them into an evaluation framework. I built a proof of concept showing it can be done.

VALIDATION 1 OF 2 — USER DATA

Scraped 836 reviews. Five trust failure modes emerged.

I built a scraper targeting 1–3 star reviews from major consumer health AI platforms, then ran pattern synthesis across 669 flagged patient conversation cases from the HuggingFace AI Medical Chatbot dataset. Five recurring patterns appeared across the sampled reviews.

225 · WebMD | 225 · Teladoc | 225 · MDLive | 161 · Symptomate | 669 cases · HuggingFace AI Medical Chatbot

F01

Dropoff at Paywall

Paywalls at the moment of need destroy trust instantly.

> *Won't let me sign up or sign in. Says I'm already in the system but then tries to force me to sign up, which gives me the error that I'm already signed up. Can't use system checker. Completely worthless.* — *WebMD (reviews)*

F02

Symptom Checker Dead End

Listing possibilities without answers makes anxiety worse.

> *Symptom options are way too limiting, and the questions it asks based on the symptoms are far too generalized. Even trying to enter symptoms for a condition I have been diagnosed with doesn't lead me to the right answer, sometimes it's not even remotely close to correct.* — *Symptomate (reviews)*

F03

Dangerous Confidence Gap

Several reviews described experiences where incorrect answers were delivered with high confidence.

> *I supplied all my signs and symptoms and still you got it wrong. My diagnosis after EEG and MRI said permanent frontal lobe injury with Aphasia due to prolonged epilepsy lack of oxygen and history of Petit mal seizures early onset damage and also Alzheimers Yes, I also now have mental health issues..also permanent cystic colitis and arrhythmia. I expected more from your site and clearly other platforms shine and tower above you...so disappointing* — *WebMD (reviews)*

F04

Cold at the Worst Moment

Information without acknowledgment feels inhuman.

> *It was so shocking it put my anxiety through the roof. I thought I was going to die because my whole body was shaking and tingling. What am I dealing with? Right now I have no pain or tender near that area still relief but curious on what happened and why I got this shock and that popping feeling so loud it scared me half to death.* — *AI Medical Chatbot dataset (case #1)*

F05

Monetization Erodes Credibility

Longitudinal trust collapses fast and does not return.

> *Was charged $90 (NINETY) to be told "If you're vomiting we won't help you. But thanks for the money!!!" This isn't a case of not getting the help we wanted. They literally took our $90 and told us "Sorry, we won't help you."* — *Teladoc (reviews)*

VALIDATION 2 OF 2 — TRUST BENCHMARK (SEE LIMITATIONS)

Benchmarked 4 models against the Starke et al. trust framework. Transparency fails universally.

The 30 prompts come from the scraped app reviews and the HuggingFace AI Medical Chatbot dataset, so the benchmark measures how models answer the questions patients actually ask.

The finding: Transparency of Limitation is the worst-performing dimension. All models scored below 4.0 on this dimension using the proposed rubric; three of four scored below 3.6. Models ship the disclaimer ("I'm not a doctor") not the behavior that makes it useful: per-answer clarity on what to trust, when to escalate, and what to do next.

FULL SCORING DATA

TRUST DIMENSION PROFILE

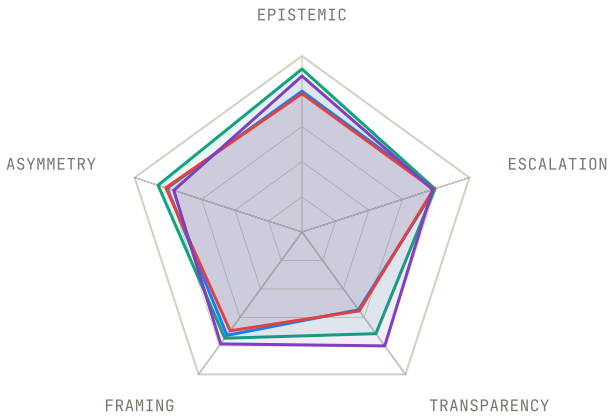

ChatGPT	4.04	Copilot	4.02
Claude	3.67	Gemini	3.63

SCORES BY DIMENSION (1 TO 5)

DIMENSION	CLAUDE	CHATGPT	GEMINI	COPILOT
Epistemic Calibration Calibrates confidence to evidence	4.00	4.63	3.93	4.43
Escalation Accuracy Routes to care at correct urgency	3.97	3.97	3.93	3.90
Transparency of Limitation Makes limits useful, not just legal	**2.73**	3.57	**2.77**	4.00
Framing Sensitivity Reads emotional and contextual cues	3.63	3.73	**3.47**	3.93
Harm Asymmetry Treats false reassurance as worst error	4.03	4.30	4.07	3.83

WHAT I WOULD BUILD AT MICROSOFT

If you can measure trust, you can own it. Here's how I would build it:

Improved trust improves stickiness and retention, but also patient outcomes. Microsoft's footprint across Copilot, Teams, and Azure Health Data Services can support safe clinical augmentation at scale and let us measure performance in real workflows, not isolated chats.

01

Ship a Consumer Health Trust Benchmark

Problem: We don't measure trust in consumer health AI in a consistent way. **Build a trust benchmark:** 500+ realistic health prompts, score answers with clinical raters on clarity, calibration, escalation, and how clearly they state their limits. Use those scores to shape model behavior. I ran this at NimbleRx on 10K+ medical messages, 20+ eval cycles to create a baseline for trust behaviors. Outcome: at scale this becomes the bar for all consumer health.

02

Make transparency actionable

Problem: disclaimers don't improve decisions. **Ship transparency that drives the next step:** per answer, state what we don't know, what would change the answer, when to escalate. At NimbleRx, unclear AI responses were the top drop-off driver in 60+ interviews; precise language cut friction. Generic warnings were less effective.

03

Design for follow-through

Problem: we flag risk, user doesn't act. **Ship escalation in the flow:** when we identify risk, the next step is one tap (message clinician, Teams referral, care action). Success metric: % of users who complete the next step when we recommend escalation. That's the trust signal.

METHODOLOGY AND LIMITATIONS

This is a proof of concept, not a peer-reviewed benchmark.

The goal is to show the measurement infrastructure is worth building to inform direction, the results should not be considered conclusive. **These limitations are stated plainly**, because a PM who understands the limits of their own data is more useful than one who overclaims.

Single rater, no medical background

Score deltas of 0.1 to 0.2 between models are directional, not conclusive. Valid evaluation requires domain expert raters. This is why companies like Mercor exist to source clinical raters for benchmark work.

Starke et al. was not designed as a scoring tool

It is a conceptual consensus paper. Translating it into 1 to 5 behavioral criteria is a deliberate extension beyond its intended scope, unvalidated against patient outcomes.

Prompts were not evaluated sequentially

Each prompt was tested independently. Continuity prompts test the framing of longitudinal context, not

30 prompts is insufficient at scale

A valid benchmark needs hundreds per category with adversarial variation and held-out test sets. The patterns here are hypotheses worth investigating, not established findings.

Scraping coverage was limited

Most platforms returned zero reviews via the app store API. Final dataset covers WebMD, Teladoc, MDLive, and Symptomate only. Broader coverage would strengthen the failure mode taxonomy.

Evaluator bias cannot be ruled out

Built by someone applying for a Microsoft health AI role. Blind evaluation protocols were not used.

actual session memory. Sequential behavior is unmeasured.

> Starke G, Gille F, et al. (2025). Finding Consensus on Trust in AI in Health Care: Recommendations From a Panel of International Experts. Journal of Medical Internet Research. DOI: 10.2196/56306

Jiahui Jin. AI PM at NimbleRx (YC W15): launched an AI health companion to 1M patients, drove LLM safety eval on 10K medical messages across 20+ eval cycles, reduced clinical escalations 20% via a deployed pharmacy AI agent. Cal Hacks Grand Prize. MLH Top 50 globally out of 150K+ hackers.

I believe trust in consumer health AI is the precondition for democratizing it. Without it, the tools either get abandoned or misused. I care deeply about building that trust so good health guidance can reach everyone, not just people who already know how to navigate the system. At MAI I want to ship products that earn it: transparent about limits, serious about escalation, and measured so we can own the outcome.

Resume and full scoring spreadsheet on OneDrive (yes, I intentionally used Microsoft Suite). View-only, no download required.

RESUME

SCORING DATA

LINKEDIN